

Assignment Project Exam Help

Architectures & P <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Dr Nguyen Tran
School of Computer Science

Previously...

- Basic definition of a distributed system
 - “*A collection of independent computers that appears to its users as a single coherent system.*”

Assignment Project Exam Help

- Real-world examples for distributed systems

<https://eduassistpro.github.io/>

- Transparency helps the users obtain a single coherent system

Add WeChat edu_assist_pro

Outline - Architectures

- Software Architectures
- System Architectures
 - ~~Centralized architectures~~ Assignment Project Exam Help
 - The Client-Side
 - The Layered<https://eduassistpro.github.io/>
 - ~~Decentralized architectures~~ Add WeChat edu_assist_pro
 - The Peer-to-Peer Organization
 - ~~Hybrid architectures~~
 - Edge server systems
 - Collaborative distributed systems

Software Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Software Architectures

Logical organization of components in distributed systems

- **Component:** A modular unit with well-defined interfaces that is replaceable
- **Connector:** a mechanism that mediates communication, coordination among components.

Layer-based Architectures vs. Object-based Architectures

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- In (a), requests (resp. responses) go downward (resp. upward)
- In (b), objects communicate through Remote Procedure Calls (RPCs)

Software Architectures

Event-based Architectures

vs.

Data-centered Architectures

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- (a) Event-based architecture: communication through events, that **optionally carry data** (subscribers get their desired events delivered)
- (b) Data-centered architecture: through a shared repository, that **contains data** (e.g., files in a distributed file system, web-based distributed system, Twitter)

Centralized Architectures

Architectures & Processes

Week 2, COMP3221

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Centralized Architectures

The basic client-server model

- The client/server requests/provides the services
- The client and the server can be hosted on different machines.

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- The communication follows a request-reply model.
- Examples: ??

The Client-Server Model

Stateless vs Stateful server



Client



Stateless
server

Assignment Project Exam Help

?

<https://eduassistpro.github.io/>

Hi, I'm comp1, may I have th f file 5?

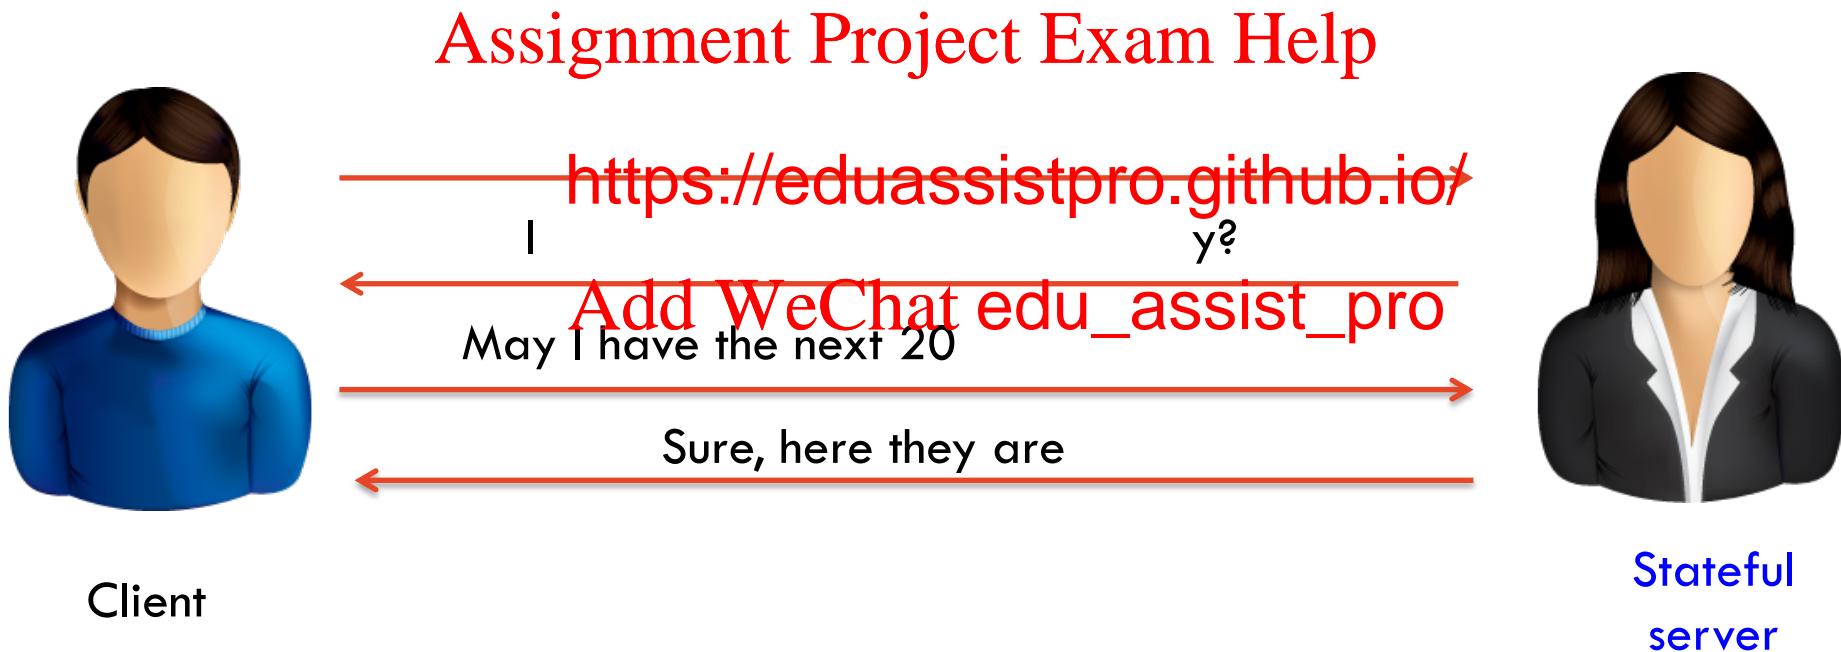
~~Add WeChat edu_assist_pro~~

Sure, you have the credentials, attached are the lines

←

The Client-Server Model

Stateless vs Stateful server



The Client-Server Model

Stateless vs Stateful server

- Stateless server: does not record the state of its clients
- Stateful server: maintains persistent information about its clients
(client->file)

<https://eduassistpro.github.io/>

	Stateless server	Stateful server
State	No info kept	Persistent
Request	Self-contained	Can be split, generally faster
Upon failure	No recovery needed	State recovery needed (explicit deletion)
Example	Network file system (NFSv3)	Andrew file system (AFS)

The Layered Organization

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Application layering

Traditional three-layered view:

1. The user interface layer

- contains the feature to control the application

Assignment Project Exam Help

2. The processing

<https://eduassistpro.github.io/>

- contains the function of the application

3. The data layer

- contains the data of the application

Application layering (cont'd)

Example: a search engine request spanning the traditional three layers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Multi-tiered architectures

Physical two-tiered architecture



Assignment Project Exam Help



<https://eduassistpro.github.io/>

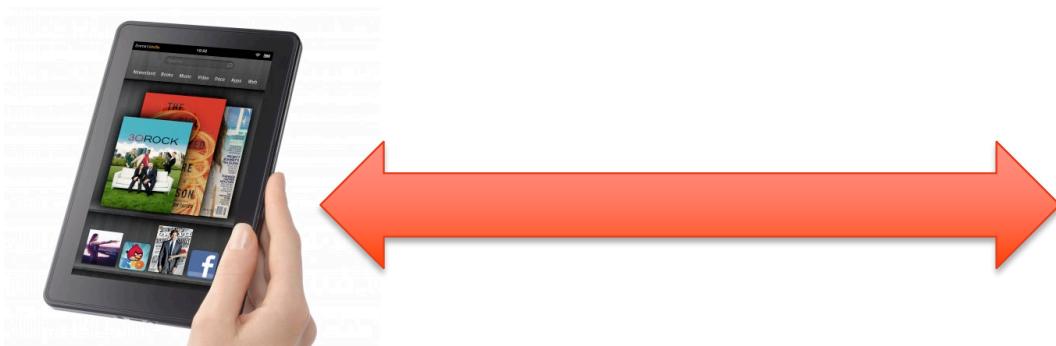
Add WeChat edu_assist_pro

Multi-tiered architectures (cont'd)

A single machine can act both as a client and a server

Example: Cloud computing

- *Cloud computing:* the delivery of computation or storage as a service to end-users.
- The servers handle most of the computation upon request and sends back the results
 - One server asks
 - Another does the computation
 - ...



Decentralized Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

The peer-to-peer model

Every machine acts similarly

- Every machine is both a client and a server
- No centralized control: the responsibility is distributed evenly
Assignment Project Exam Help
- Even the program is similar
- Forms an **overlay**
<https://eduassistpro.github.io/>
 - nodes are processors,
Add WeChat edu_assist_pro
 - links possible communication channels.

Structured Peer-to-Peer Model

The overlay network is constructed using a deterministic procedure

- **Chord** is an example of a **Distributed Hash Table (DHT)**

As a node: Assignment Project Exam Help

- I have a success <https://eduassistpro.github.io/>
- I have a predec
- I have some ~~shortcuts~~ Add WeChat edu_assist_pro
nodes to speedup delivery of
requests
- I am responsible of a subset of
the system data items (based on
my unique identifier)

Hybrid Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Edge-Server Systems

- Servers are placed at the edge of the network.
- Edge servers serve content, possibly after applying filtering and transcoding functions
- One edge server acts as a origin server for all content, while taking advantage of mized delivery of content, e.g. thro<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Collaborative Distributed Systems

- How do you first get started ?
 - Often a traditional client-server models is used.
 - Once joined fully decentralized schemes can be used.

Assignment Project Exam Help

- Example: BitTorr

<https://eduassistpro.github.io/>

- 35% of world-wide internet traffic
- Used for Linux distribution, software movies
- Goal: quickly replicate large files to large number of clients





Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Web server hosts a .torrent file (th, hash, tracker's URL...)
Add WeChat edu_assist_pro
- A tracker (server or a DHT) tracks downloaders/owners of a file
- Files are divided into chunks (256kb-1MB)
- Downloaders download chunks from themselves (and owners)
- **Tit-for-tat:** the more one shares (**server**), the faster it can download (**client**)

Conclusion

- Client and server are used to identify the role of communication participant
- Client and server roles may run on:
 - The same machine <https://eduassistpro.github.io/>
 - Distinct machines with very different IP addresses
 - Distinct machines with similar IP addresses
- Most modern distributed systems are hybrid architectures

Assignment Project Exam Help

Architecture & Processes
Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Dr. Nguyen Tran
School of Computer Science

Introduction

- **Previously:** a distributed system gives the illusion of a single system to a user thanks to transparency

Assignment Project Exam Help

- **Now:** a single system provides multiple resources to multiple users through <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Outline - Processes

- A Very Brief History

- UNIX Processes and Threads

Assignment Project Exam Help

- Multi-threading

<https://eduassistpro.github.io/>

- Multi-threaded Server

Add WeChat edu_assist_pro

A Very Brief History

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Evolution of Operating Systems (OSes)

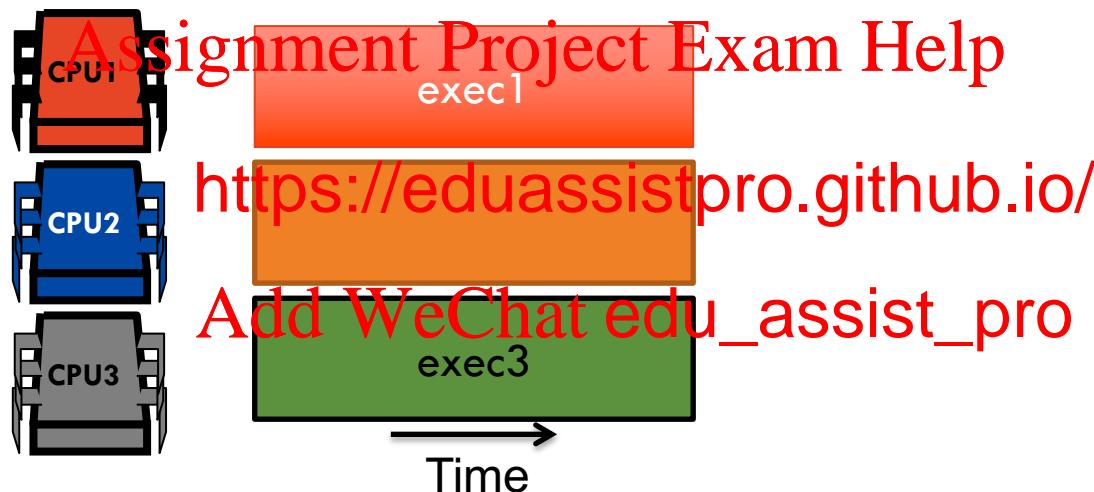
- The hardware evolved
 - Hardware was expensive, humans were cheap (e.g., Multics 1960's)
 - Hardware was cheaper, humans were expensive (e.g., desktop computer 1980's) **Assignment Project Exam Help**
 - Hardware are very cheap, humans are very expensive (e.g., handheld devices 2000's) <https://eduassistpro.github.io/>
- The OS interaction had to adapt
 - Batch: one execution at a time
 - Multiprogramming: multiple program executions simultaneously
 - Timeshare: split time into slots allocated for different program execution
 - GUI: multiple interfaces to access a system from different end-points
 - Ubiquitous devices: each user possesses her own computational device

Uniprogramming vs. Multiprogramming

- **Uniprogramming:** one program execution at a time
 - MS/DOS
 - early Macintosh
 - Batch processing
 - No longer acceptable!
- **Multiprogramming:** many programs execution at a time
 - Multics, UNIX/Linux, Windows NT,
- Many program executions on personal computers:
 - Browsing the web
 - Sending emails
 - Typing a letter
 - Burning a DVD

Concurrency

- Assume a single central processing unit (CPU)
- **Problem:** How to give the illusion to the users of multiple CPUs?

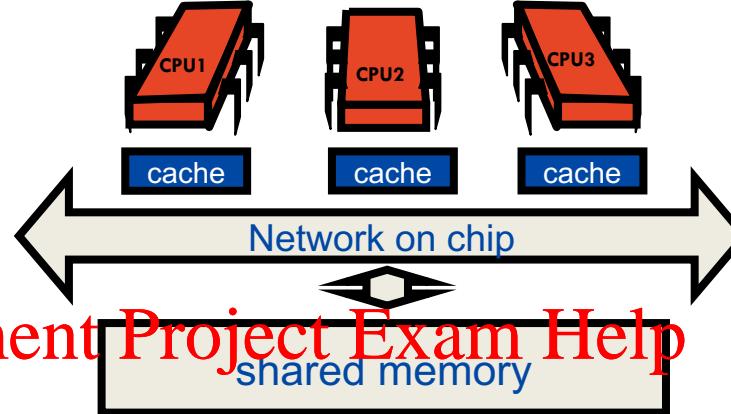


- **Solution:** **scheduling** program executions, one after the other during small fractions of the time



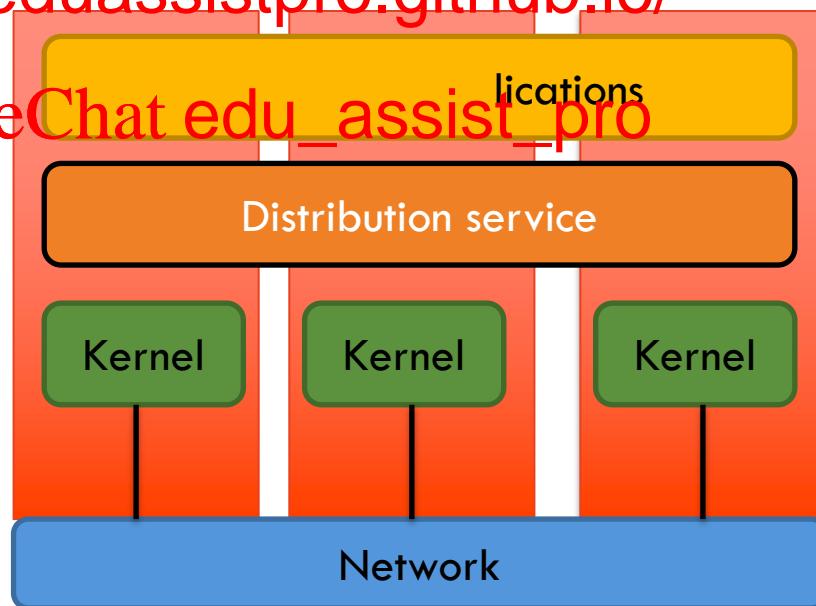
Concurrency

- Multi/many-cores on a single machine



<https://eduassistpro.github.io/>

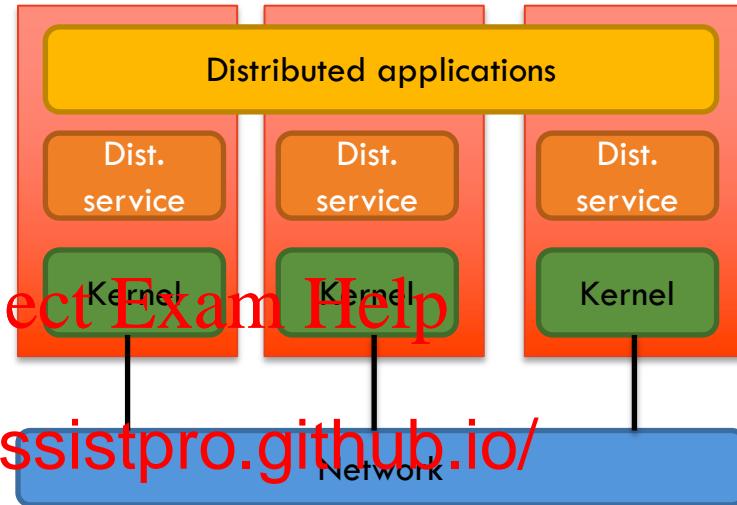
- Distributed operating systems
 - This is a single system image, the system maintains a single copy of the resources



Concurrency

- Network operating systems

- Machines provide resources to other machines (e.g. UNIX rlogin)

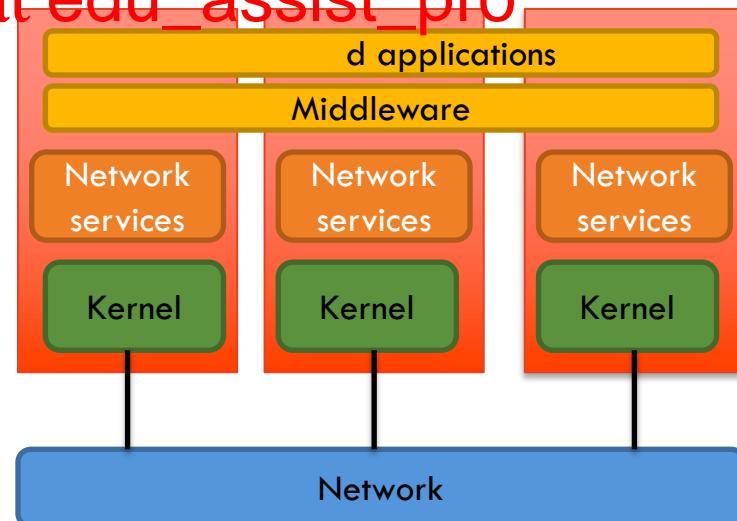


<https://eduassistpro.github.io/>

- Middleware

- A layer over the network services providing general services to applications in a very transparent manner (systems can differ)

Add WeChat `edu_assist_pro`



Processes and Threads

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Processes

Process: defined as a program in execution

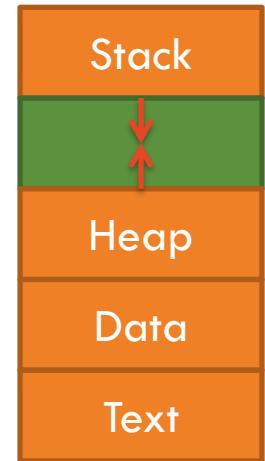
- Operating system creates a virtual CPU for each process
Assignment Project Exam Help
- OS makes sure that processes do not interfere with each other. This is done by giving each process its own virtual memory space.
inadvertently affect <https://eduassistpro.github.io/>, i.e. Secure !
- This concurrent transparency comes at a high cost:
 - For every process (virtual CPU), OS must create a completely independent address space.

Address space

Address space: A unit of management of a process's virtual memory

- Process address space:
 - **Stack**: temporary data extensible towards lower virtual addresses
 - **Heap**: memory virtual addresses
 - **Data** section: global variable
 - **Text** region: program code

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro



Process

Process: an abstraction representing a program execution

- A process switches from a state to another:

- When there is timer interrupt goes off, explicit yield, I/O, etc.

Assignment Project Exam Help

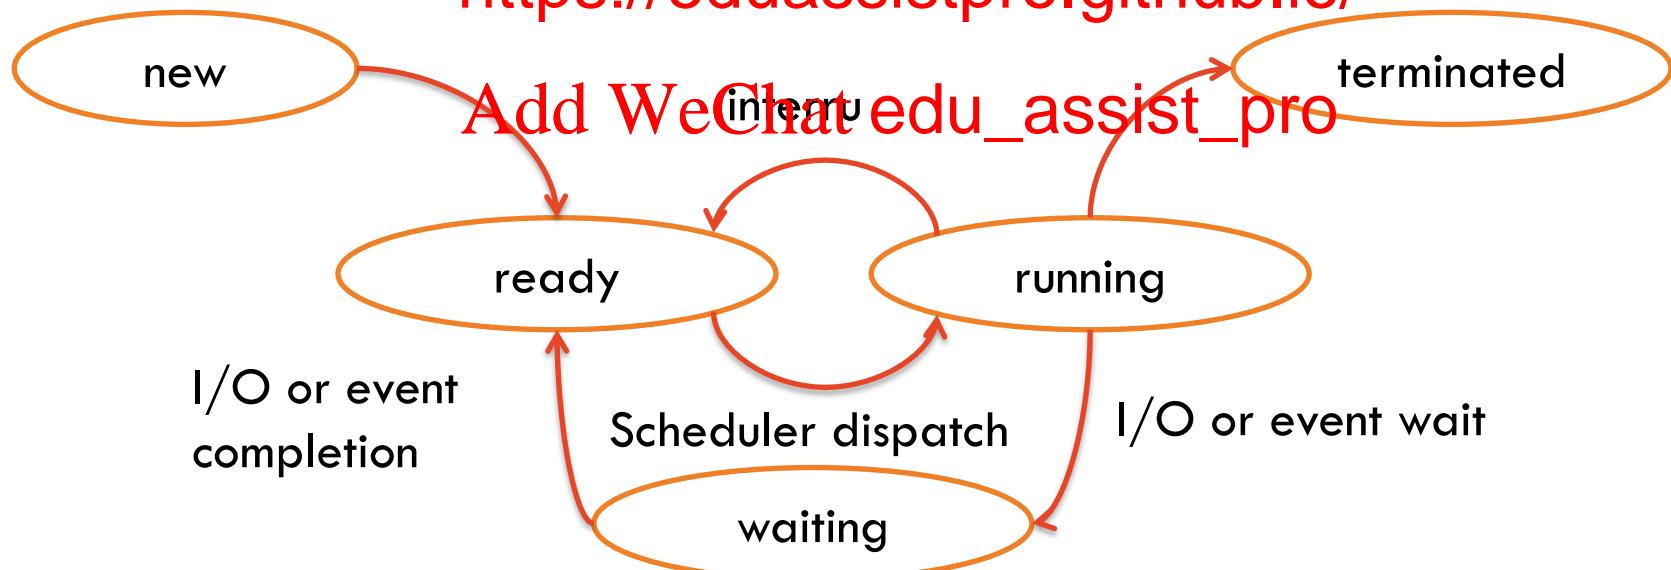
exec1

exec1.1

exec2.1

Time

<https://eduassistpro.github.io/>



Process management

Process creation:

- Process allocation: the act of choosing the host of the process
 - The system provides command to execute a process at an idle workstation
 - The system chooses processors to execute it
- Execution environment: <https://eduassistpro.github.io/> initialized content & open files
 - Static: program text, heap and stack regions created from a list
 - Dynamic: UNIX fork shares program text and copies stack and heap regions

Process management



- Switching from one process to another has become one of the most expensive operations due to **Context-Switches**
 - CPU context: re
 - Modify register stack pointer, etc.
 - Invalidate address translation cache (Translation lookaside buffer)

Inter-Process Communication (IPC):

- Pipes/Message queues/Shared memory segments
- Requires costly context switches

Context-switch

IPC requires kernel intervention

- S1. Process A switches from user to kernel mode
 - Changing the memory map in the MMU (memory management unit).
 - Flushing the TLB (translation lookaside buffer)
- S2. Process cont (switches from A to B)
- S3. Process B switches from kernel to user mode
 - Again requires changing the MMU
 - Again requires flushing the TLB

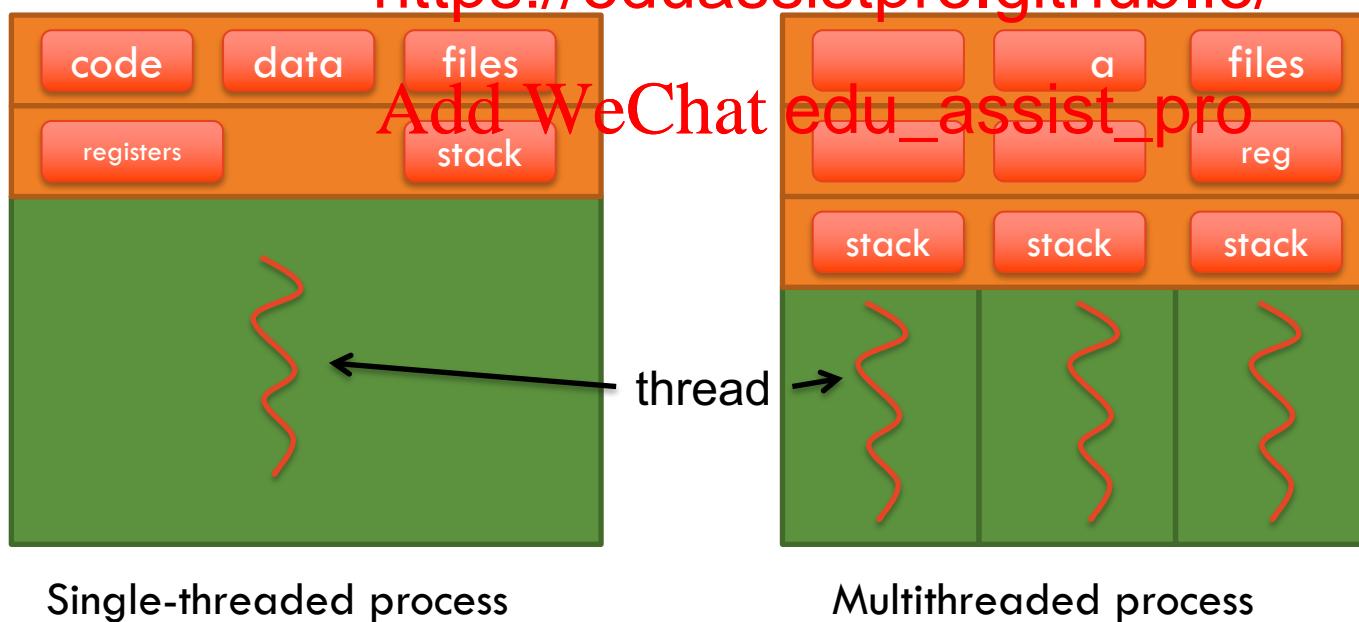
Assignment Project Exam Help

Add WeChat edu_assist_pro

Thread

Threads: smallest unit of CPU utilization

- Multiple threads per process
- Portable OS Interface (POSIX) Threads
 - Current activity: program counter
 - Stack: temporary data
 - No data, code, data, files, registers, stack



Thread

- Portable OS Interface (POSIX) Threads
 - Current activity: program counter
 - Stack: temporary data
 - No data, code (with other threads)
<https://eduassistpro.github.io/>
- Communication between threads
 - Does not need IPC
 - No context switches required (when purely at user-level)

Thread Implementation

User-level threads library vs. kernel-scheduled thread

- **User-level thread library:**
 1. Cheap to create and destroy
 2. Blocking system call freezes the entire process of the thread (and other threads)
 3. Cheap context-sw
 - Few instructions
 - No need to change memory map in MMU
 - No need to flush the TLB
- **Kernel-scheduled threads:**
 1. Costly to create and destroy
 2. Do not block the current process (and other threads) upon blocking system calls (I/O)
 3. Costly context-switch (similar to process context switch)
 - Needs to change memory map in MMU
 - Needs to flush the TLB

Processes vs. threads

- Processes
 - Isolated: prevents one process from interfering with another
 - Inefficient: starting/terminating a process and context switches are costly
- Threads
 - Non-isolated: avoiding incorrectness makes programming harder
 - Efficient: a thread is a lightweight version of a process
 - Multiple threads of control per process

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

nces makes

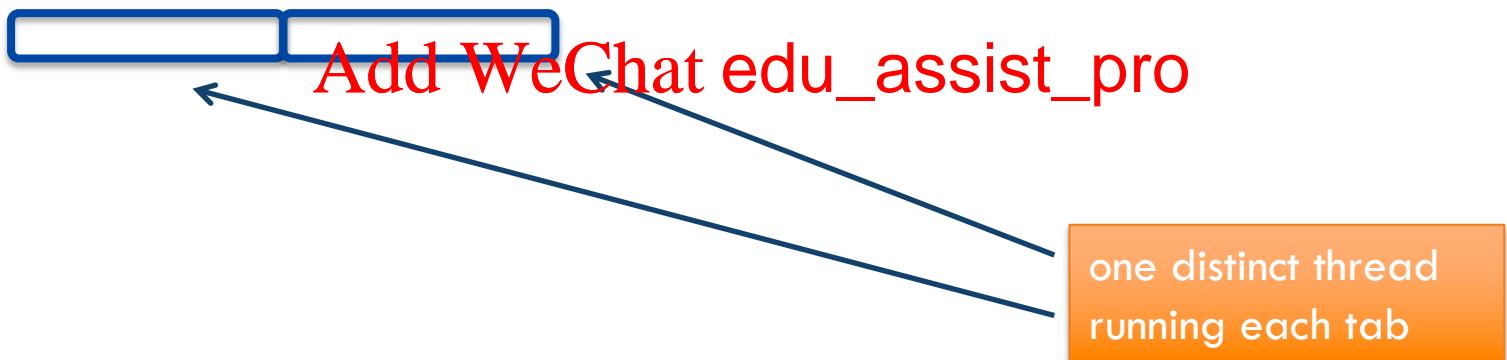
Client multi-threading

Example: The Google Chrome web browser

- Multithreading: each tab runs its own thread
- A JavaScript error does not crash the main Chrome process:
Only one tab freezes, the user can close it independently of others
- Takes benefit of m d runs on a separate core



<https://eduassistpro.github.io/>



Conclusion

- Concurrency has been used for decades since hardware was powerful enough to address multiple human needs

Assignment Project Exam Help

- Processes are ~~heavy~~ threads are lightweight but ~~not~~ <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Threads require synchronization between each other

What's Next ?

- **Tutorial on Wednesday.**
 - Multithreading
- Read ChaptAssignment Project Exam Help
- Next week: How <https://eduassistpro.github.io/> tributed systems ?
Add WeChat edu_assist_pro
- See you all next week !