

Assignment Project Exam Help

Communication (<https://eduassistpro.github.io/>)

Add WeChat edu_assist_pro

Dr Nguyen Tran

School of Computer Science

Previously...

- Previous lecture:
 - Message-based communication is complex (e.g., routing towards destination, subject to message losses)

Assignment Project Exam Help

- Today's lecture:
 - How to avoid m <https://eduassistpro.github.io/>
 - How to give the pens locally (not remotely)? [Add WeChat edu_assist_pro](#)
 - One to many communication

Outline

- The Problem of Message Loss

- The TCP/IP Solution

Assignment Project Exam Help

- Multicast commu

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Problem of Message Loss

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Message loss

Cause

- Networks are in general **unreliable**
- Messages can be **lost** (never been delivered even if sent)
- Examples:
 - A server receives too many requests and cannot treat all
 - A router drops the message because its queue is full

Message losses may impact the computation of a distributed system

Message loss

Coordinated Attack Problem



- Constraints of the problem
 - Two armies, each positioned on one of the mountains surrounding a battlefield (distance is large)
 - Can only communicate via messenger (messenger is a third party)
 - Messengers can be killed before reaching the battlefield (message losses)
- Goal: they want to coordinate an attack
 - If they attack at different times, they both die
 - If they attack at the same time, they win

Message loss

Coordinated Attack Problem (con't)

- There is no protocols to make sure they will win!



12h!

Assignment Project Exam Help

time

<https://eduassistpro.github.io/>

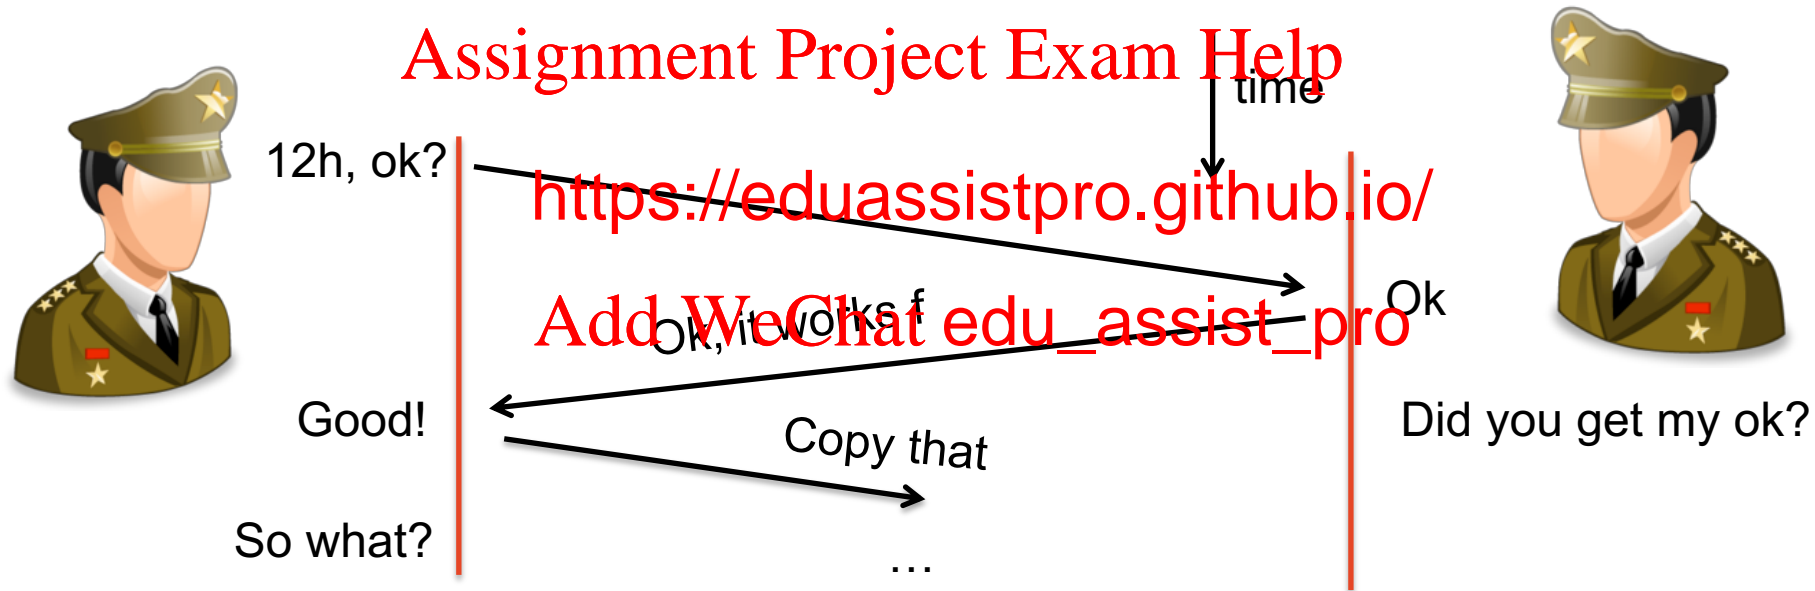
Add WeChat edu_assist_pro



Message loss

Coordinated Attack Problem (con't)

- There is no protocols to make sure they will win!



Message loss

Analogy in networking

- Constraints of the problem
 - Two remote entities of a distributed system
 - Can only communicate through messages
 - The network is unreliable
- Goal: they want to make sure to get simultaneously

This is impossible, even if all messages go through

TCP/IP

Communication 2/2
Week 4, COMP3221

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

TCP: Overview RFCs: 793,1122,1323, 2018, 2581

- Point-to-point:
 - one sender, one receiver
- Connection-oriented: Project Exam Help
- Full duplex data
 - bi-direction <https://eduassistpro.github.io/>
- Flow controlled: Add WeChat edu_assist_pro
 - sender will not overwhelm receiver
- Congestion controlled:
 - TCP congestion and flow control set window size

TCP seq. numbers, ACKs

sequence numbers:

- byte stream “number” of first byte in segment’s data

acknowledgement

- seq # of next byte expected from other side

- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

window size
N



sequence number space

sent
ACKed

sent, not-
yet ACKed
 (“in-
flight”)

usable
but not
yet sent

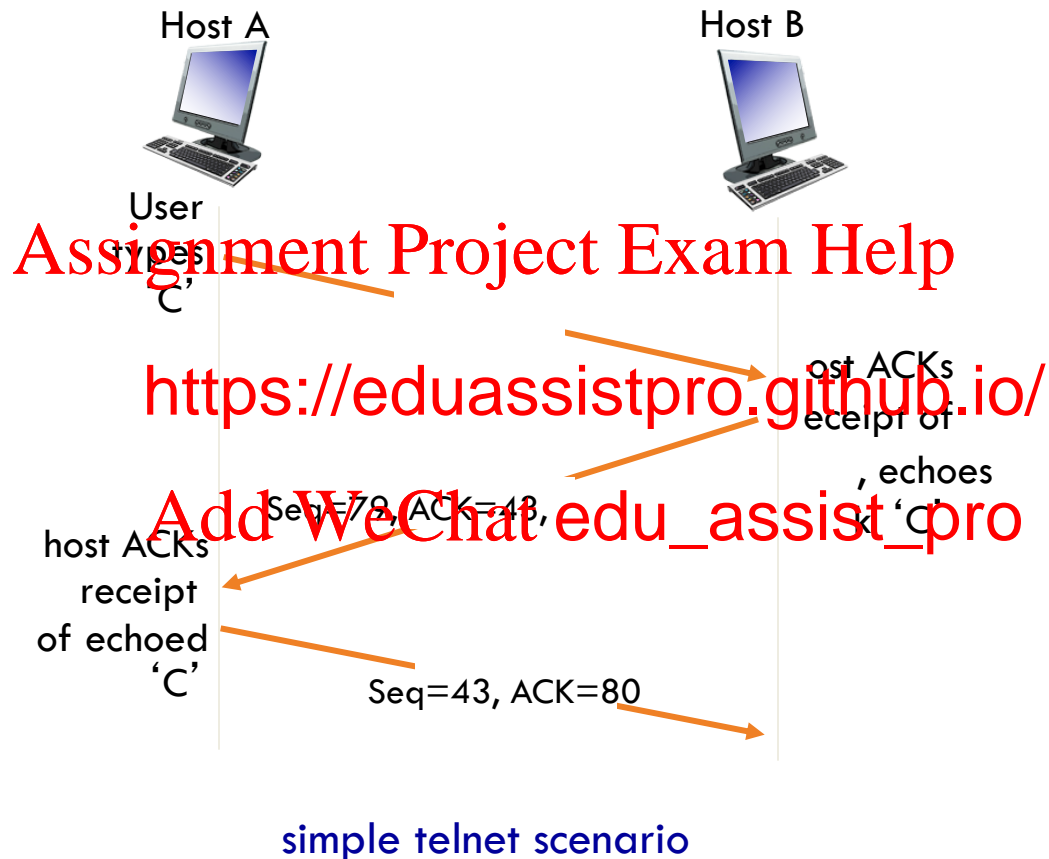
not
usable

incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

A

TCP seq. numbers, ACKs



TCP sender events:

data rcvd from app:

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not running
 - think of timer as for oldest unacked segment
 - expiration interval: `TimeoutInterval`

timeout:

- retransmit segment that caused timeout

- restart timer

ack rcvd:

- update what is known to be ACKed
- start timer if there are still unacked segments

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

TCP round trip time, timeout

Q: how to set TCP timeout value?

- longer than RTT
 - but RTT varies
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss

Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
 - retransmissions
- RTT will vary, want RTT “smoother”
 - average several recent measurements, not just current **SampleRTT**

Assignment Project Exam Help


<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.25$



TCP round trip time, timeout

- **timeout interval:** **EstimatedRTT** plus “safety margin”
 - large variation in **EstimatedRTT** -> larger safety margin
- estimate SampleRTT deviation from EstimatedRTT:

DevRTT = (

https://eduassistpro.github.io/

(typically,

Add WeChat edu_assist_pro

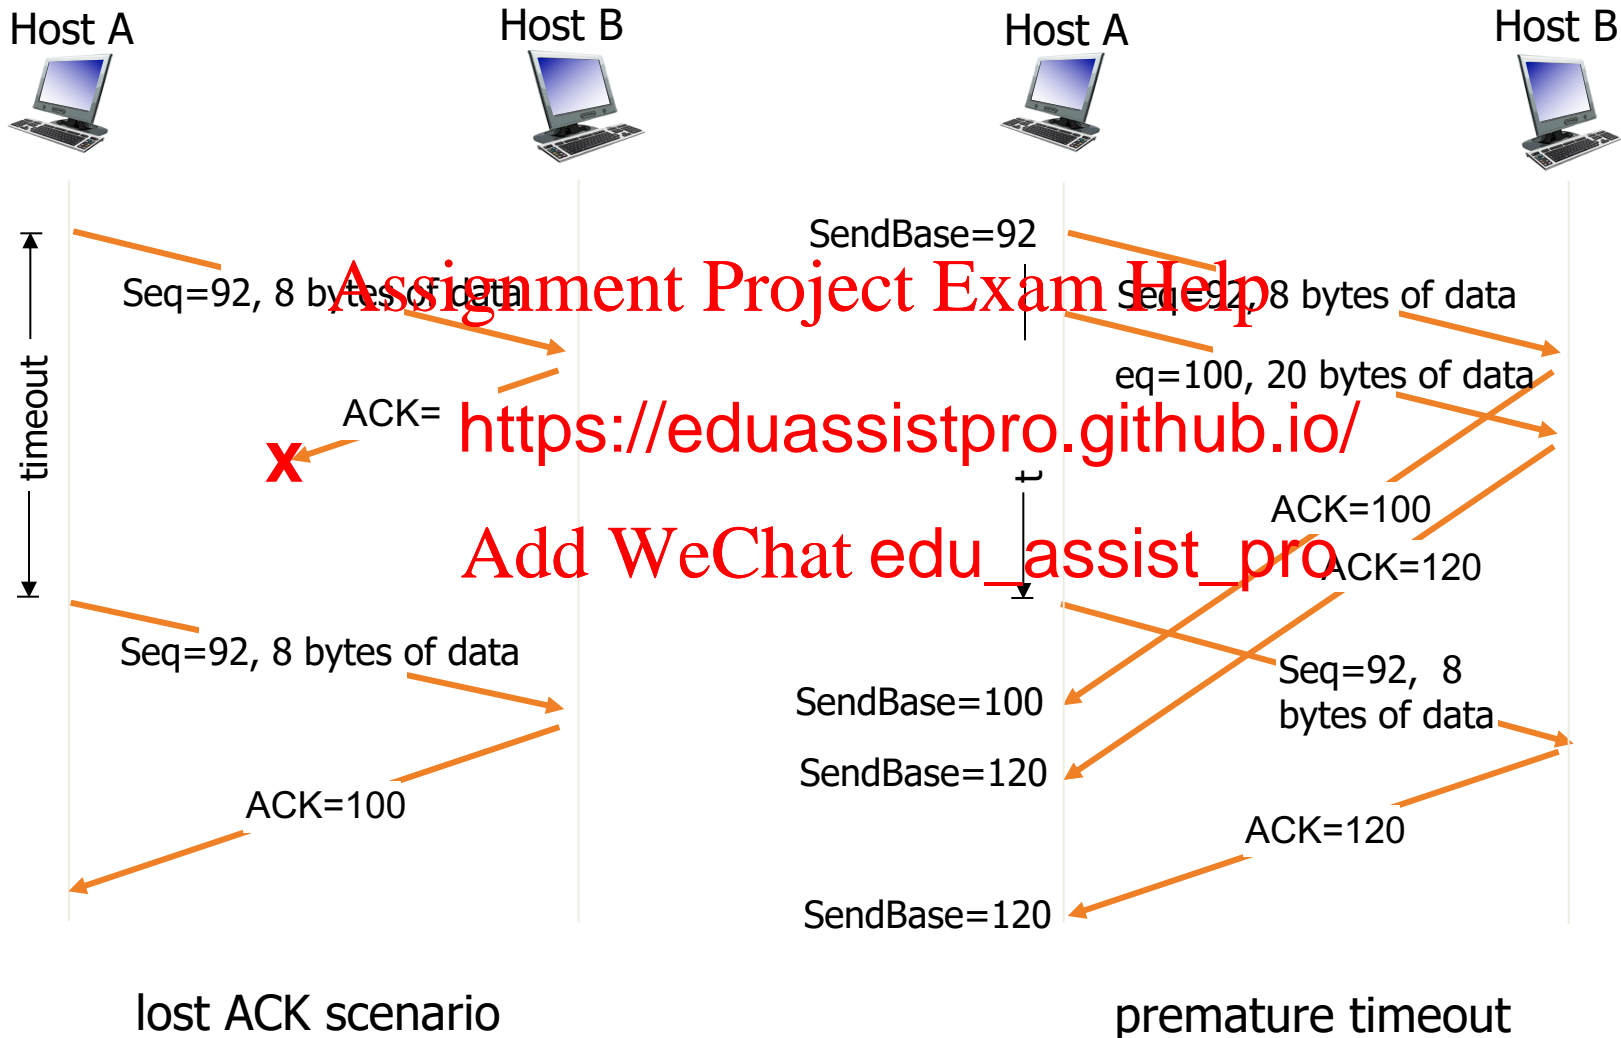
TimeoutInterval = EstimatedRTT + 4*DevRTT



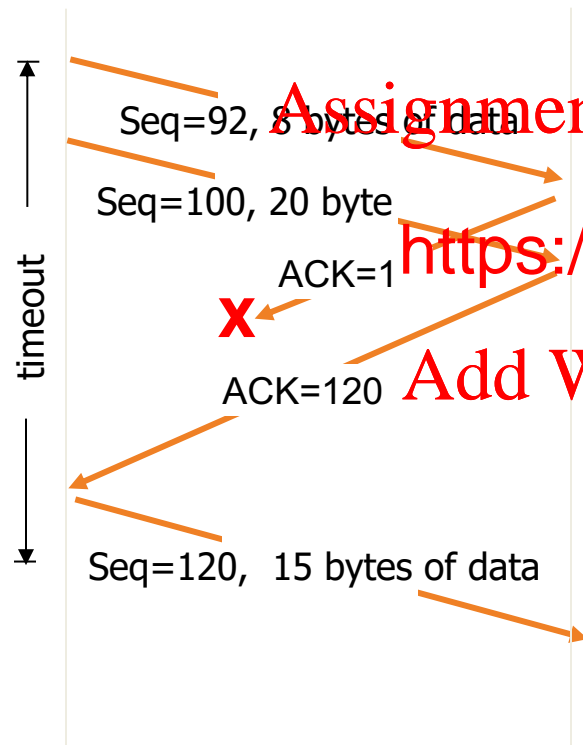
↑
estimated RTT

↑
“safety margin”

TCP: retransmission scenarios



TCP: retransmission scenarios



cumulative ACK

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

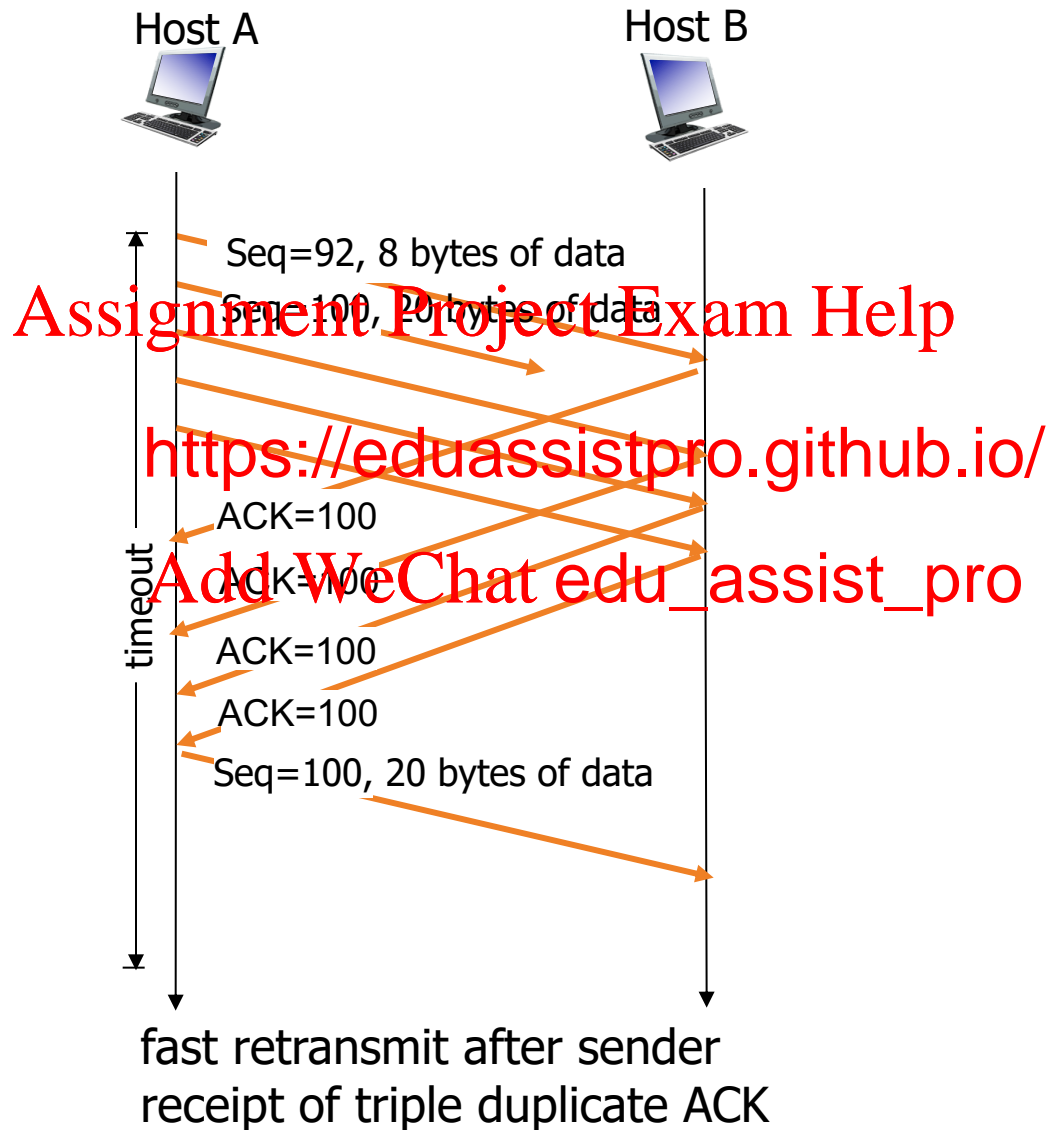
TCP fast retransmit

- time-out period often relatively long:
 - long delay before resending lost packet
- detect lost segm duplicate ACKs.
 - sender often sends many segments back-to-back
 - if segment is lost, there will likely be many duplicate ACKs.

TCP fast retransmit

if sender receives 3
ACKs for same data
(duplicate ACKs”),
unacked
with smallest
■ likely that unacked
segment lost, so don't
wait for timeout

TCP fast retransmit



TCP flow control

application may
remove data from
TCP socket buffers

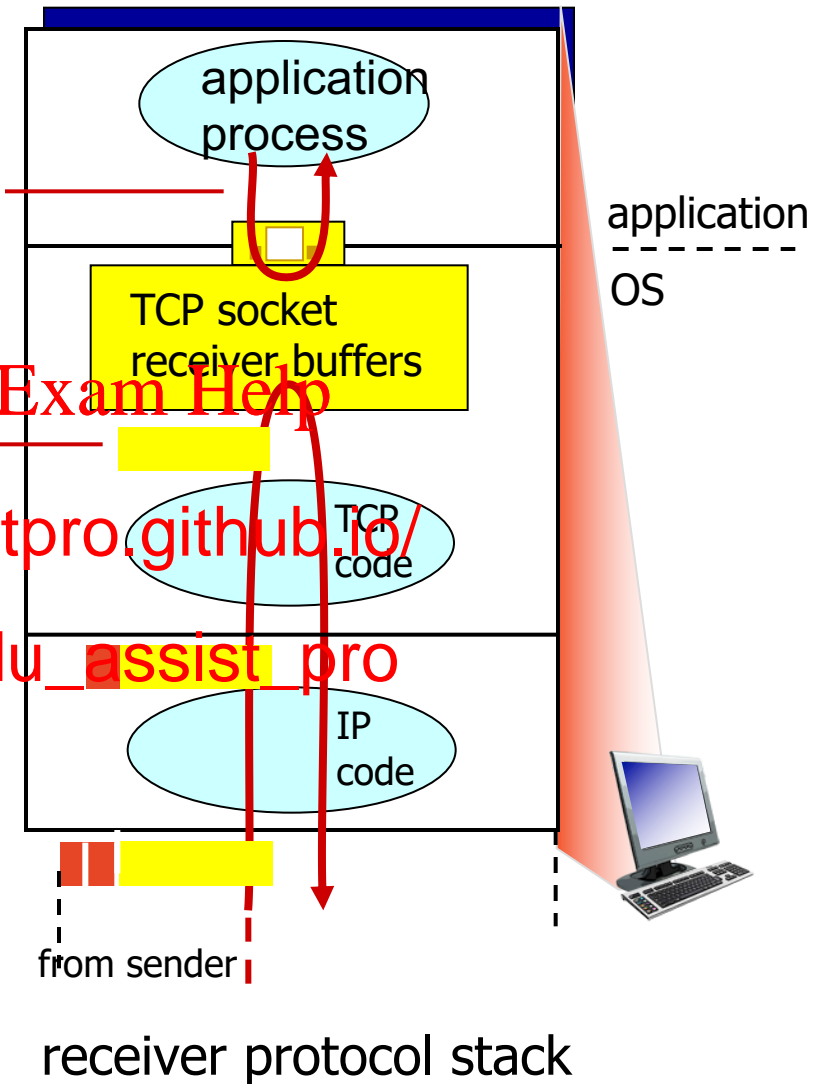
Assignment Project Exam Help
... slower than TCP
receiver is delivering

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

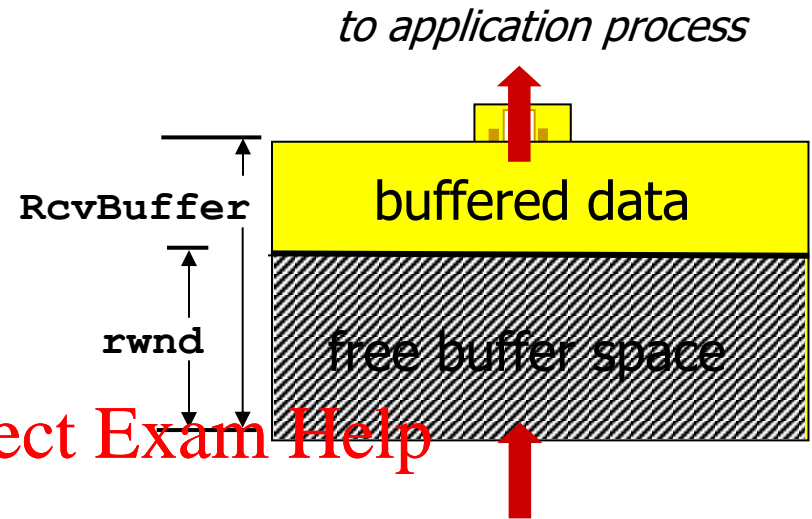
flow control

receiver controls sender, so
sender won't overflow
receiver's buffer by transmitting
too much, too fast



TCP flow control

- receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments



- **RcvBuffer** is a socket option (typically 4096 bytes)
- many operating systems automatically adjust **RcvBuffer**

- sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- guarantees receive buffer will not overflow

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

TCP segment payloads

receiver-side buffering

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

Principles of congestion control

congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from f
- manifestations:
 - lost packets (buffer overflows)
 - long delays (queueing in router buffers)
- a top-10 problem!


TCP congestion control

- *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase **cwnd** by 1 MSS (Maximum Segment Size) every RTT until loss detected
 - *multiplicative decrease*: **cwnd** half after loss

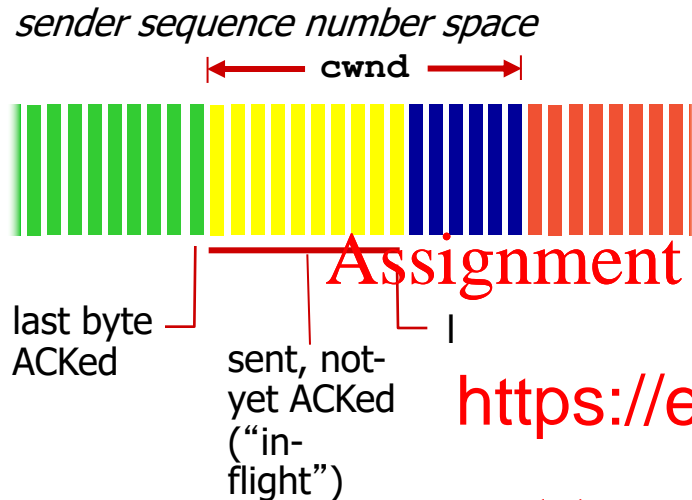
<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](https://eduassistpro.github.io/)

AIMD saw tooth behavior: probing for bandwidth



TCP Congestion Control: details



TCP sending rate:

- roughly: send cwnd bytes, wait RTT for ACKS, then send more bytes

<https://eduassistpro.github.io/>

- sender limits transmission:

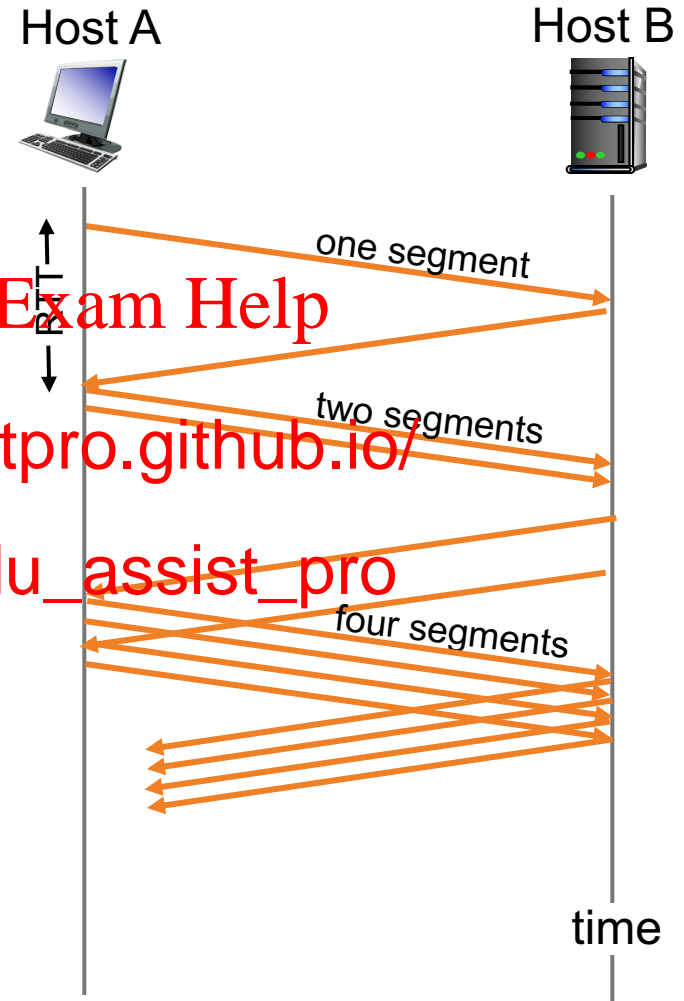
$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- Congestion window (**cwnd**) is dynamic function of perceived network congestion

$$\text{Add WeChat edu_assist_pro} \quad \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

TCP Slow Start

- when connection begins, increase rate exponentially until first loss event:
 - initially $cwnd = 1 \text{ MSS}$ (Maximum Segment Size)
 - double $cwnd$
 - done by incrementing $cwnd$ for every ACK received
- summary: initial rate is slow but ramps up exponentially fast



TCP: detecting, reacting to loss

- Loss indicated by timeout:
 - **cwnd** set to 1 MSS;
 - window then grows exponentially (as in slow start) to threshold, then grows linearly
- Loss indicated by duplicate ACKs (indicating some segments received out of order)
 - **cwnd** is cut in half
 - window then grows linearly
- TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate acks)

TCP: switching from slow start to Congestion Avoidance

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to $1/2$ of its value before timeout.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to $1/2$ of **cwnd** just before loss event

Connection Management


before exchanging data, sender/receiver “**handshake**”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters, e.g. MSS, rwnd, etc.
- connection-oriented


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

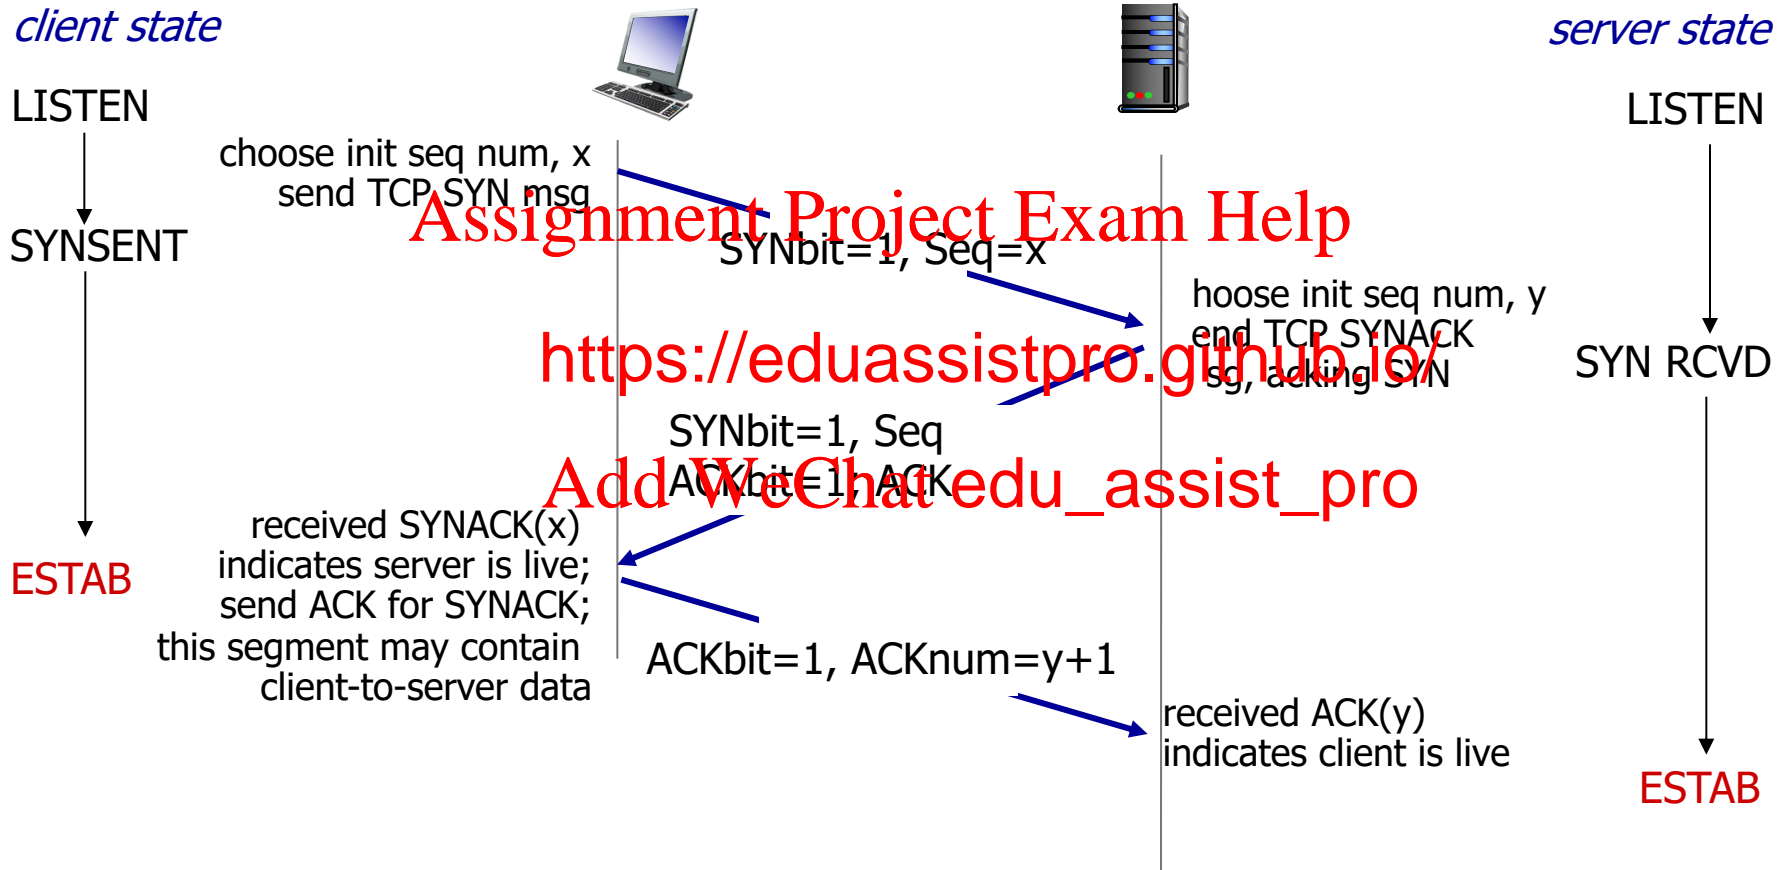


```
Socket clientSocket =  
    newSocket("hostname", "port  
    number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

TCP 3-way handshake



TCP: closing a connection

client state

ESTAB

`clientSocket.close()`

FIN_WAIT_1

can no longer
send but can
receive data

FINbit=1, seq=x

FIN_WAIT_2

wait for

can still
send data

TIMED_WAIT

timed wait
for $2 \times \text{max}$
segment lifetime

FINbit=1,

can no longer
send data

ACKbit=1; ACKnum=y+1

CLOSED

server state

ESTAB

CLOSE_WAIT

LAST_ACK

CLOSED

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

HTTP Overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - **client**: browser that receives, (using, and "displays"
 - **server**: Web server sends (using HTTP protocol) objects in response to requests



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

HTTP overview

Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains no information about past client requests

Assignment Project Exam Help

<https://eduassistpro.github.io/>

TTPS ?

Add WeChat edu_assist_pro

ver TLS

ort Layer

Security), e.g. over SSL

Related terminology

- **Broadcast: one-to-all communication**
 - Action of sending a message to all nodes of the system
 - Typically used for relatively small systems, like IP broadcast as part of Ethernet in LANs
- **Unicast: one-to-one communication**
 - In contrast with broadcast, action of distribution to a single receiver
 - Used generally in streaming applications
- **Anycast: one-to-random-one communication**
 - Send a message to an IP address range and obtain a response from any potential receiver, whose IP address belongs to this range
 - Used with UDP and TCP
- **Multicast: one-to-many communication**
 - Action of sending a message to multiple nodes of the system (not necessarily all nodes)
 - This term is used in many contexts (network, algorithm)

Conclusion (Transportation layer)

- Message losses can have dramatic consequences
- TCP/IP protocol suite hides these losses from the application level
- Socket, RPC, use
 - Sockets are complex
 - RPC is more transparent for the client
- Multicast communication scales well for large distributed systems

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro